

## ▼ Configuration

```
#Parameters
PROJECT_NAME = 'ML1010_Weekly'
ENABLE_COLAB = True

#Root Machine Learning Directory. Projects appear underneath
GOOGLE_DRIVE_MOUNT = '/content/gdrive'
COLAB_ROOT_DIR = GOOGLE_DRIVE_MOUNT + '/MyDrive/Colab Notebooks'
COLAB_INIT_DIR = COLAB_ROOT_DIR + '/utility_files'

LOCAL_ROOT_DIR = '/home/magni/Documents/ML_Projects'
LOCAL_INIT_DIR = LOCAL_ROOT_DIR + '/utility_files'
```

## ▼ Bootstrap Environment

```
#add in support for utility file directory and importing
import sys
import os

if ENABLE_COLAB:
    #Need access to drive
    from google.colab import drive
    drive.mount(GOOGLE_DRIVE_MOUNT, force_remount=True)

    #add in utility directory to syspath to import
    INIT_DIR = COLAB_INIT_DIR
    sys.path.append(os.path.abspath(INIT_DIR))

    #Config environment variables
    ROOT_DIR = COLAB_ROOT_DIR

else:
    #add in utility directory to syspath to import
    INIT_DIR = LOCAL_INIT_DIR
    sys.path.append(os.path.abspath(INIT_DIR))

    #Config environment variables
    ROOT_DIR = LOCAL_ROOT_DIR

#Import Utility Support
from jarvis import Jarvis
jarvis = Jarvis(ROOT_DIR, PROJECT_NAME)

import mv python utils as mvutils
```

```
Mounted at /content/gdrive
Wha...where am I?
I am awake now.
```

```
I have set your current working directory to /content/gdrive/MyDrive/Colab Notebooks/ML1
The current time is 11:57
Hello sir. Extra caffeine may help.
```

## ▼ Setup Runtime Environment

```
if ENABLE_COLAB:
    #!pip install scipy -q
    #!pip install scikit-learn -q
    #!pip install pycaret -q
    #!pip install matplotlib -q
    #!pip install joblib -q
    #!pip install pandasql -q

    display('Google Colab enabled')
else:
    display('Google Colab not enabled')

#Common imports
import json
import gzip
import pandas as pd
import numpy as np
import matplotlib
import re
import nltk
import matplotlib.pyplot as plt

pd.set_option('mode.chained_assignment', None)
nltk.download('stopwords')
%matplotlib inline

import warnings

warnings.filterwarnings("ignore")

'Google Colab enabled'
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

## ▼ Load Data

```
jarvis.showAllDataFiles()
```

Here are all your available data files

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data [Empty directory]
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/Jarvis
```

```
---[ gz][ pk1]--> 02_NLP_ReviewTextData.pkl.gz (30.08 MB)
```

```
---[ gz][ pk1]--> 02_NLP_SummaryData.pkl.gz (2.88 MB)
```

```
---[ gz][ csv]--> movie_reviews_cleaned.csv.gz (14.73 MB)
```

```
---[ gz][ csv]--> wk3_task_data.csv.gz (33.47 KB)
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/Jarvis/01_original [Empty directory]
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/Jarvis/02_working [Empty directory]
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/Jarvis/03_train
```

```
---[ gz][ pk1]--> 02_NLP_SummaryData.pkl.gz (2.88 MB)
```

```
---[ gz][ pk1]--> 02_NLP_TitleData.pkl.gz (1.43 MB)
```

```
---[ gz][ pk1]--> 03_NLP_ReviewTextData.pkl.gz (10.91 MB)
```

```
---[ gz][ pk1]--> 03_NLP_SummaryData.pkl.gz (1.62 MB)
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/Jarvis/04_test
```

```
---[ gz][ csv]--> pima-indians-diabetes.csv.gz (8.53 KB)
```

```
---[ gz][ csv]--> wk3_task_data.csv.gz (33.47 KB)
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/ML1010-Group-Project [Empty directory]
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/ML1010-Group-Project/01_original
```

```
---[ gz][ json]--> Cell_Phones_and_Accessories_5.json.gz (161.24 MB)
```

```
---[ gz][ json]--> meta_Cell_Phones_and_Accessories.json.gz (343.33 MB)
```

```
[D] /content/gdrive/MyDrive/Colab Notebooks/data/ML1010-Group-Project/02_working
```

```
[*][ pk1]-----> 01_Cellphone_small.pkl (45.46 MB)
```

```
---[ gz][ pk1]--> 01_NLP_ReviewText_Narrow_1.pkl.gz (6.88 MB)
```

```
---[ gz][ pk1]--> 01_NLP_ReviewText_Narrow_2.pkl.gz (170.55 MB)
```

```
---[ gz][ pk1]--> 01_NLP_ReviewText_Narrow_3.pkl.gz (295.59 MB)
```

```
[*][ pk1]-----> 01_NLP_ReviewText_small.pkl (28.94 MB)
```

```
[*][ pk1]-----> 01_NLP_Summary_small.pkl (3.82 MB)
```

```
[*][ pk1]-----> 01_NLP_Title_small.pkl (2.73 MB)
```

```
---[ gz][ pk1]--> 01_NL_ReviewText_All(new).pkl.gz (593.23 MB)
```

```
---[ gz][ pk1]--> 01_NL_ReviewText_All.pkl.gz (592.92 MB)
```

```
---[ gz][ pk1]--> 01_NL_ReviewText_textSplit.pkl.gz (15.78 MB)
```

```
[*][ pk1]-----> 02_Cellphone.pkl (46.32 MB)
```

```
[*][ pk1]-----> 02_NLP_ReviewTextData.pkl (87.00 MB)
```

```
[*][ pk1]-----> 02_NLP_SummaryData.pkl (8.32 MB)
```

```
[*][ pk1]-----> 02_NLP_TitleData.pkl (16.71 MB)
```

```
[*][ pk1]-----> 03_Cellphone.pkl (46.31 MB)
```

```
[*][ pk1]-----> 03_NLP_ReviewTextData.pkl (28.94 MB)
```

```
[*][ pk1]-----> 03_NLP_ReviewText_Narrow.pkl (17.13 MB)
```

```
[*][ pk1]-----> 03_NLP_SummaryData.pkl (3.82 MB)
```

```
[*][  pk1]-----> 03_NLP_TitleData.pkl (2.73 MB)
[*][  pk1]-----> 04_NLP_ReviewText_Narrow.pkl (16.95 MB)
[*][  pk1]-----> 05_NLP_ReviewText_Narrow.pkl (66.15 MB)
[*][  pk1]-----> 05_NLP_ReviewText_Narrow_full.pkl (207.91 MB)

[D] /content/gdrive/MyDrive/Colab Notebooks/data/ML1010-Group-Project/03_train [Empty]
[D] /content/gdrive/MyDrive/Colab Notebooks/data/ML1010-Group-Project/04_test [Empty]
```

```
df = pd.read_csv(jarvis.DATA_DIR + '/complaints.csv.gz')
```

```
mvutils.exploreDataframe(df, numRecords=1)
#df.info(verbose=True)
```

```
dataframe shape: (2355756, 18)
```

```
dataframe info:
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 2355756 entries, 0 to 2355755
```

```
Data columns (total 18 columns):
```

#	Column	Dtype
0	Date received	object
1	Product	object
2	Sub-product	object
3	Issue	object
4	Sub-issue	object
5	Consumer complaint narrative	object
6	Company public response	object
7	Company	object
8	State	object
9	ZIP code	object

```
#delete all null values for narrative
```

```
df = df[pd.notnull(df['Consumer complaint narrative'])]
```


```
df.reset_index(inplace=True, drop=True)
```

```
mvutils.exploreDataframe(df)
```

```
dataframe shape: (821617, 18)

dataframe info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 821617 entries, 0 to 821616
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Date received                         821617 non-null object
1   Product                              821617 non-null object
2   Sub-product                          769445 non-null object
3   Issue                                821617 non-null object
4   Sub-issue                            651600 non-null object
5   Consumer complaint narrative          821617 non-null object
6   Company public response               398420 non-null object
7   Company                              821617 non-null object
8   State                                817979 non-null object
9   ZIP code                             643495 non-null object
10  Tags                                  131213 non-null object
11  Consumer consent provided?            821617 non-null object
12  Submitted via                         821617 non-null object
13  Date sent to company                  821617 non-null object
14  Company response to consumer          821616 non-null object
15  Timely response?                      821617 non-null object
16  Consumer disputed?                    164062 non-null object
17  Complaint ID                           821617 non-null int64
```

```
mvutils.displayClassBalance(df, 'Product')
```



Bottom 1 in dataframe

```
df = df.groupby('Product', group_keys=False).apply(lambda x: x.sample(frac=0.01))

received Product product Issue issue complaint public

import importlib
importlib.reload(mvutils)

<module 'mv_python_utils' from '/content/gdrive/MyDrive/Colab Notebooks/utility_files/mv
```

```

col = ['Product', 'Consumer complaint narrative']
df = df[col]
df.columns

Index(['Product', 'Consumer complaint narrative'], dtype='object')

df.columns = ['Product', 'Consumer_complaint_narrative']

df['category_id'] = df['Product'].factorize()[0]
from io import StringIO
category_id_df = df[['Product', 'category_id']].drop_duplicates().sort_values('category_id')
category_to_id = dict(category_id_df.values)
id_to_category = dict(category_id_df[['category_id', 'Product']].values)

df.head()


```

	Product	Consumer_complaint_narrative	category_id
<b>593793</b>	Bank account or service	Purchased a Jumbo C.D. from CIT Bank XXXX/XXXX...	0
<b>256852</b>	Bank account or service	I had applied for a CitiGold Checking account ...	0
<b>658921</b>	Bank account or service	RE Case number XXXX On XX/XX/2016, I hired XXX...	0
-----	Bank account or	This is my second complaint against Chase	-

```

import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,6))
df.groupby('Product').Consumer_complaint_narrative.count().plot.bar(ylim=0)
plt.show()

```



```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
tfidf = TfidfVectorizer(sublinear_tf=True, min_df=5, norm='l2', encoding='latin-1', ngram_range=(1, 2))
```

```
features = tfidf.fit_transform(df.Consumer_complaint_narrative).toarray()
```

```
labels = df.category_id
```

```
features.shape
```

```
(8215, 22919)
```

```
edi
```

```
from sklearn.feature_selection import chi2
```

```
import numpy as np
```

```
N = 2
```

```
for Product, category_id in sorted(category_to_id.items()):
```

```
    features_chi2 = chi2(features, labels == category_id)
```

```
    indices = np.argsort(features_chi2[0])
```

```
    feature_names = np.array(tfidf.get_feature_names())[indices]
```

```
    unigrams = [v for v in feature_names if len(v.split(' ')) == 1]
```

```
    bigrams = [v for v in feature_names if len(v.split(' ')) == 2]
```

```
    print("# '{}':".format(Product))
```

```
    print("    . Most correlated unigrams:\n        . {}".format('\n        . '.join(unigrams[-N:])))
```

```
    print("    . Most correlated bigrams:\n        . {}".format('\n        . '.join(bigrams[-N:])))
```

```
    # 'Bank account or service':
```

```
        . Most correlated unigrams:
```

```
            . citigold
```

```
            . overdraft
```

```
        . Most correlated bigrams:
```

```
            . citigold checking
```



```

    . overdraft fees
# 'Checking or savings account':
    . Most correlated unigrams:
        . deposit
        . bank
    . Most correlated bigrams:
        . debit card
        . checking account
# 'Consumer Loan':
    . Most correlated unigrams:
        . car
        . santander
    . Most correlated bigrams:
        . motor finance
        . months loan
# 'Credit card':
    . Most correlated unigrams:
        . amex
        . card
    . Most correlated bigrams:
        . macy credit
        . credit card
# 'Credit card or prepaid card':
    . Most correlated unigrams:
        . merchant
        . card
    . Most correlated bigrams:
        . american express
        . credit card
# 'Credit reporting':
    . Most correlated unigrams:
        . experian
        . equifax
    . Most correlated bigrams:
        . disputed equifax
        . equifax ignored
# 'Credit reporting, credit repair services, or other personal consumer reports':
    . Most correlated unigrams:
        . inquiries
        . report
    . Most correlated bigrams:
        . identity theft
        . credit report
# 'Debt collection':
    . Most correlated unigrams:
        . collection
        . debt
    . Most correlated bigrams:
        . debt collection
        . collect debt
# 'Money transfer, virtual currency, or money service':
    . Most correlated unigrams:

```

```

from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer

```

```

from sklearn.naive_bayes import MultinomialNB

X_train, X_test, y_train, y_test = train_test_split(df['Consumer_complaint_narrative'], df['P
count_vect = CountVectorizer()
X_train_counts = count_vect.fit_transform(X_train)
tfidf_transformer = TfidfTransformer()
X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts)

clf = MultinomialNB().fit(X_train_tfidf, y_train)

print(clf.predict(count_vect.transform(["This company refuses to provide me verification and

['Credit reporting, credit repair services, or other personal consumer reports']

#I trimmed data and used a subset so it doesn't find narrative in the dataframe
df[df['Consumer_complaint_narrative'] == "This company refuses to provide me verification and

```

Product	Consumer_complaint_narrative	category_id
---------	------------------------------	-------------

```
df.head()
```

	Product	Consumer_complaint_narrative	category_id
<b>593793</b>	Bank account or service	Purchased a Jumbo C.D. from CIT Bank XXXX/XXXX...	0
<b>256852</b>	Bank account or service	I had applied for a CitiGold Checking account ...	0
<b>658921</b>	Bank account or service	RE Case number XXXX On XX/XX/2016, I hired XXX...	0
-----	Bank account or	This is my second complaint against Chase	-

```

print(clf.predict(count_vect.transform(["I am disputing the inaccurate information the Chex-S

['Credit reporting, credit repair services, or other personal consumer reports']

```

```

#I trimmed data and used a subset so it doesn't find narrative in the dataframe
df[df['Consumer_complaint_narrative'] == "I am disputing the inaccurate information the Chex-

```

Product	Consumer_complaint_narrative	category_id
---------	------------------------------	-------------

```

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import LinearSVC

```

```

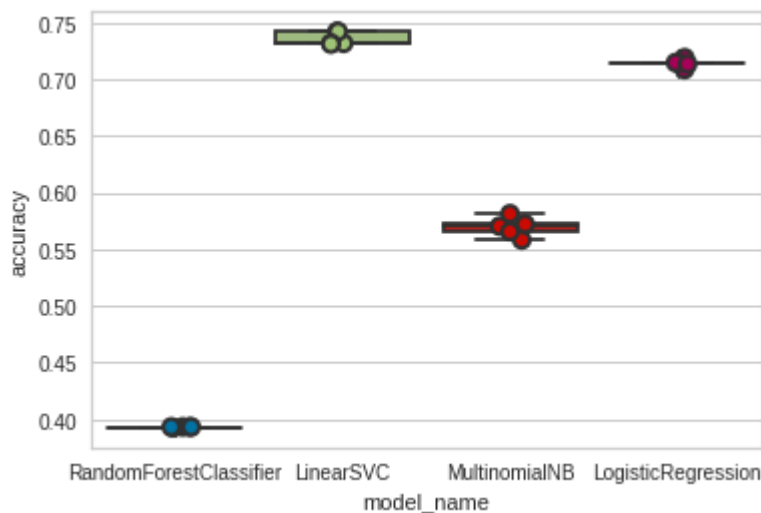
from sklearn.model_selection import cross_val_score

models = [
    RandomForestClassifier(n_estimators=200, max_depth=3, random_state=0),
    LinearSVC(),
    MultinomialNB(),
    LogisticRegression(random_state=0),
]
CV = 5
cv_df = pd.DataFrame(index=range(CV * len(models)))
entries = []
for model in models:
    model_name = model.__class__.__name__
    accuracies = cross_val_score(model, features, labels, scoring='accuracy', cv=CV)
    for fold_idx, accuracy in enumerate(accuracies):
        entries.append((model_name, fold_idx, accuracy))
cv_df = pd.DataFrame(entries, columns=['model_name', 'fold_idx', 'accuracy'])

import seaborn as sns
sns.boxplot(x='model_name', y='accuracy', data=cv_df)
sns.stripplot(x='model_name', y='accuracy', data=cv_df,
              size=8, jitter=True, edgecolor="gray", linewidth=2)

plt.show()

```



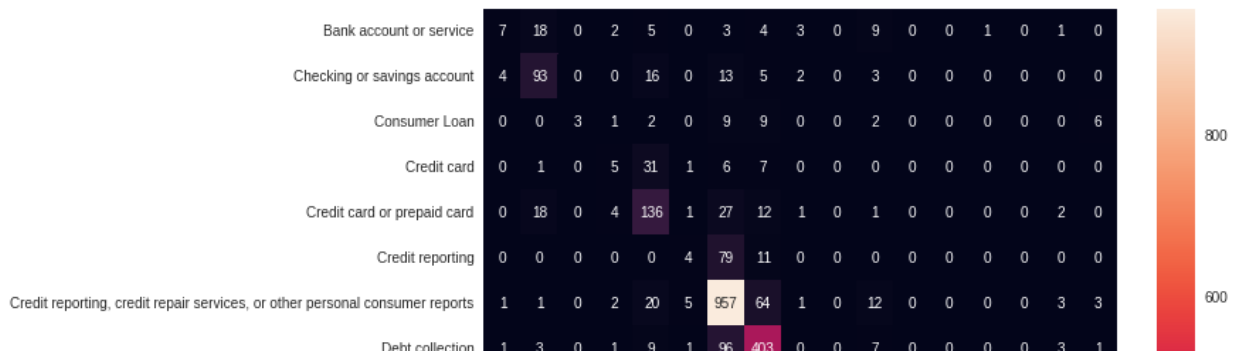
```
cv_df.groupby('model_name').accuracy.mean()
```

```

model_name
LinearSVC          0.735606
LogisticRegression 0.713938
MultinomialNB      0.569446
RandomForestClassifier 0.393061
Name: accuracy, dtype: float64

```

```
model = LinearSVC()
X_train, X_test, y_train, y_test, indices_train, indices_test = train_test_split(features, la
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
from sklearn.metrics import confusion_matrix
conf_mat = confusion_matrix(y_test, y_pred)
fig, ax = plt.subplots(figsize=(10,10))
sns.heatmap(conf_mat, annot=True, fmt='d',
             xticklabels=category_id_df.Product.values, yticklabels=category_id_df.Product.val
plt.ylabel('Actual')
plt.xlabel('Predicted')
plt.show()
```



```
category_id_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 17 entries, 593793 to 286016
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Product         17 non-null    object
1   category_id     17 non-null    int64
dtypes: int64(1), object(1)
memory usage: 408.0+ bytes
```

```
0 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17
```

```
from IPython.display import display
```

```
for predicted in category_id_df.category_id:
    for actual in category_id_df.category_id:
        if predicted != actual and conf_mat[actual, predicted] >= 6:
            print("{}' predicted as '{}' : {} examples.".format(id_to_category[actual], id_to_category[predicted], conf_mat[actual, predicted]))
            display(df.loc[indices_test[(y_test == actual) & (y_pred == predicted)]][['Product', 'Category']])
            print('')
```

'Bank account or service' predicted as 'Checking or savings account' : 18 examples

	Product	Consumer_complaint_narrative
525836	Bank account or service	checking account was over-drafted. I had, acco...
646988	Bank account or service	Dear Sir, On Friday, XX/XX/XXXX my wallet ( CC...
563293	Bank account or service	My boyfriend and I have an account with US Ban...
659868	Bank account or service	XXXX separate Wells Fargo Banks in XXXX, Texas...
652467	Bank account or service	My checking account with PNC Bank had a low ba...
659567	Bank account or service	I got a job offer on XXXX for a administrative...
594970	Bank account or service	Citibank advertised the promotional bonus in X...
584175	Bank account or service	I 'm writing in regards to our experience with...
687637	Bank account or service	I had slightly more than {\$1700.00} in my Bank...
624464	Bank account or service	I have been a customer with Bank of America fo...
641811	Bank account or service	At approximately XXXX on XXXX XXXX I went to m...
630554	Bank account or service	I was called out of state for a family emergen...
564049	Bank account or service	I had several accounts at Wells Fargo From XX/...
618502	Bank account or service	I deposited {\$7000.00} into my checking accoun...
629807	Bank account or service	In 2011 my debit card was lost/stolen I had ju...
646152	Bank account or service	Ally Bank refused to transfer a matured CD IRA...
575978	Bank account or service	I had XXXX charges that cleared last week, the...
693403	Bank account or service	The bank honored a check that was dated a mont...

```

model.fit(features, labels)
N = 2
for Product, category_id in sorted(category_to_id.items()):
    indices = np.argsort(model.coef_[category_id])
    feature_names = np.array(tfidf.get_feature_names())[indices]
    unigrams = [v for v in reversed(feature_names) if len(v.split(' ')) == 1][:N]
    bigrams = [v for v in reversed(feature_names) if len(v.split(' ')) == 2][:N]
    print("# '{}':".format(Product))
    print(" . Top unigrams:\n      . {}".format('\n      . '.join(unigrams)))
    print(" . Top bigrams:\n      . {}".format('\n      . '.join(bigrams)))

# 'Bank account or service':
. Top unigrams:
. bank
. referring
. Top bigrams:
. met requirements
. xxxx 15
# 'Checking or savings account':

```

```

. Top unigrams:
  . bank
  . branch
. Top bigrams:
  . overdraft fee
  . debit card
# 'Consumer Loan':
. Top unigrams:
  . santander
  . car
. Top bigrams:
  . months loan
  . new car
# 'Credit card':
. Top unigrams:
  . card
  . amex
. Top bigrams:
  . credit card
  . closed xxxx
# 'Credit card or prepaid card':
. Top unigrams:
  . card
  . discover
. Top bigrams:
  . credit limit
  . use card
# 'Credit reporting':
. Top unigrams:
  . equifax
  . experian
. Top bigrams:
  . xxxx accounts
  . xxxx contract
# 'Credit reporting, credit repair services, or other personal consumer reports':
. Top unigrams:
  . experian
  . report
. Top bigrams:
  . xxxx xxxx
  . disputed xxxx
# 'Debt collection':
. Top unigrams:
  . debt
  . collection
. Top bigrams:
  . company account
  . xxxx filed
# 'Money transfer, virtual currency, or money service':
  Top unigrams:

```

```

from sklearn import metrics
print(metrics.classification_report(y_test, y_pred, target_names=df['Product'].unique()))

```

precision

Bank account or service	0.50
Checking or savings account	0.60
Consumer Loan	1.00
Credit card	0.31
Credit card or prepaid card	0.58
Credit reporting	0.29
Credit reporting, credit repair services, or other personal consumer reports	0.76
Debt collection	0.71
Money transfer, virtual currency, or money service	0.77
Money transfers	0.00
Mortgage	0.86
Other financial service	0.00
Payday loan	1.00
Payday loan, title loan, or personal loan	0.80
Prepaid card	0.00
Student loan	0.80
Vehicle loan or lease	0.52
accuracy	
macro avg	0.56
weighted avg	0.71

80777	Checking or savings account	Hi earlier this month, I paid a 5 dollar fee t...
723665	Checking or savings account	I was banking with Wellsfargo my account was c...
504580	Checking or savings account	On XXXX/XXXX/XXXX I reserved a car through XXX...
442567	Checking or savings account	My name is XXXX XXXX and actually this is abou...
220426	Checking or savings account	I get my child support from XXXX on a Direct E...
281422	Checking or savings account	Hello, today I spoke an agent via 5th 3rd bank...
258399	Checking or savings account	On, XX/XX/XXXX all my XXXX transactions where ...
235295	Checking or savings account	I opened a CD account with Marcus Bank. I rece...
306299	Checking or savings account	We have over 90 unauthorized chargers on our U...
307247	Checking or savings account	Dear CFPB, today I received a letter in the US...
63470	Checking or savings account	On XX/XX/2020 my wallet was stolen and the per...

'Credit card' predicted as 'Credit card or prepaid card' : 31 examples.

Product

Consumer\_complaint\_narrative

888485 Credit card In XX/XX/XXXX and XX/XX/XXXX my credit card...

