

Import necessary dependencies

```
In [1]: ENABLE_COLAB=False
```

```
In [2]: if ENABLE_COLAB:
        !pip install pycaret -q
        #!/pip install https://github.com/pandas-profiling/pandas-profiling/archive/
        #!/pip install matplotlib -q
        #!/pip install pandasql -q
    else:
        display('Google Colab not enabled')

'Google Colab not enabled'
```

```
In [3]: if ENABLE_COLAB:
        from pycaret.utils import enable_colab
        enable_colab()

        from google.colab import drive
        drive.mount('/content/gdrive', force_remount=True)
    else:
        display('Google Colab not enabled')

'Google Colab not enabled'
```

```
In [5]: import os
        import sys
        print("Current working directory: {}".format(os.getcwd()))
        os.chdir('/home/magni/ML_Root/project_root/utility_files')
        print("Current working directory: {}".format(os.getcwd()))
        sys.path.append('.')

Current working directory: /home/magni/ML_Root/project_root/ML1010_Weekly
Current working directory: /home/magni/ML_Root/project_root/utility_files
```

```
In [6]: import model_evaluation_utils as meu
```

```
In [7]: import pandas as pd
        import numpy as np

        #import text_normalizer as tn

        np.set_printoptions(precision=2, linewidth=80)
```

Load and normalize data

```
In [9]: dataset = pd.read_csv('/home/magni/ML_Root/project_root/data/ML1010_Weekly/mo

# take a peek at the data
print(dataset.head())
reviews = np.array(dataset['review'])
sentiments = np.array(dataset['sentiment'])

# build train and test datasets
train_reviews = reviews[:5000]
train_sentiments = sentiments[:5000]
test_reviews = reviews[5000:7000]
test_sentiments = sentiments[5000:7000]

# normalize datasets
#norm_train_reviews = tn.normalize_corpus(train_reviews)
norm_train_reviews = train_reviews
#norm_test_reviews = tn.normalize_corpus(test_reviews)
norm_test_reviews = test_reviews
```

```

                                review sentiment
0  not bother think would see movie great supspen... negative
1  careful one get mitt change way look kung fu f... positive
2  chili palmer tired movie know want success mus... negative
3  follow little know 1998 british film make budg... positive
4  dark angel cross huxley brave new world percys... positive
```

Traditional Supervised Machine Learning Models

Feature Engineering

```
In [10]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer

# build BOW features on train reviews
cv = CountVectorizer(binary=False, min_df=0.0, max_df=1.0, ngram_range=(1,2))
cv_train_features = cv.fit_transform(norm_train_reviews)
# build TFIDF features on train reviews
tv = TfidfVectorizer(use_idf=True, min_df=0.0, max_df=1.0, ngram_range=(1,2),
                    sublinear_tf=True)
tv_train_features = tv.fit_transform(norm_train_reviews)
```

```
In [11]: # transform test reviews into features
cv_test_features = cv.transform(norm_test_reviews)
tv_test_features = tv.transform(norm_test_reviews)
```

```
In [12]: print('BOW model:> Train features shape:', cv_train_features.shape, ' Test fe
print('TFIDF model:> Train features shape:', tv_train_features.shape, ' Test
```

```
BOW model:> Train features shape: (5000, 434563) Test features shape: (2000,
434563)
TFIDF model:> Train features shape: (5000, 434563) Test features shape: (200
```

0. 434563)

Model Training, Prediction and Performance Evaluation

In [13]: `from sklearn.linear_model import SGDClassifier, LogisticRegression`

```
lr = LogisticRegression(penalty='l2', max_iter=1000, C=1)
svm = SGDClassifier(loss='hinge', max_iter=1000)
```

In [14]: `# Logistic Regression model on BOW features`
`lr_bow_predictions = meu.train_predict_model(classifier=lr,`
 `train_features=cv_train_features`
 `test_features=cv_test_features,`
`meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_`
 `classes=['positive', 'negative'])`

Model Performance metrics:

```
-----
Accuracy: 0.8605
Precision: 0.8606
Recall: 0.8605
F1 Score: 0.8605
```

Model Classification report:

```
-----
              precision    recall  f1-score   support

   positive       0.85        0.86        0.86         981
   negative       0.87        0.86        0.86        1019

   accuracy                   0.86         2000
  macro avg       0.86        0.86        0.86         2000
 weighted avg     0.86        0.86        0.86         2000
```

Prediction Confusion Matrix:

```
-----
              Predicted:
              positive  negative
Actual: positive      846       135
       negative      144       875
```

In [15]: `# Logistic Regression model on TF-IDF features`
`lr_tfidf_predictions = meu.train_predict_model(classifier=lr,`
 `train_features=tv_train_featur`
 `test_features=tv_test_features`
`meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_`
 `classes=['positive', 'negative'])`

Model Performance metrics:

```
-----
Accuracy: 0.866
Precision: 0.8661
Recall: 0.866
```

F1 Score: 0.866

Model Classification report:

	precision	recall	f1-score	support
positive	0.87	0.85	0.86	981
negative	0.86	0.88	0.87	1019
accuracy			0.87	2000
macro avg	0.87	0.87	0.87	2000
weighted avg	0.87	0.87	0.87	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	838	143

In [16]:

```
svm_bow_predictions = meu.train_predict_model(classifier=svm,
                                              train_features=cv_train_features,
                                              test_features=cv_test_features,
                                              meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_
                                              classes=['positive', 'negative']))
```

Model Performance metrics:

 Accuracy: 0.8525
 Precision: 0.8525
 Recall: 0.8525
 F1 Score: 0.8525

Model Classification report:

	precision	recall	f1-score	support
positive	0.85	0.85	0.85	981
negative	0.85	0.86	0.86	1019
accuracy			0.85	2000
macro avg	0.85	0.85	0.85	2000
weighted avg	0.85	0.85	0.85	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	829	152
negative	143	876

In [17]:

```
svm_tfidf_predictions = meu.train_predict_model(classifier=svm,
                                              train_features=tv_train_featu
                                              test_features=tv_test_feature
                                              meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_
                                              classes=['positive', 'negative']))
```

Model Performance metrics:

```
-----
Accuracy: 0.881
Precision: 0.881
Recall: 0.881
F1 Score: 0.881
```

Model Classification report:

```
-----
              precision    recall  f1-score   support

   positive       0.88       0.87       0.88        981
   negative       0.88       0.89       0.88       1019

   accuracy                0.88        2000
  macro avg       0.88       0.88       0.88        2000
 weighted avg     0.88       0.88       0.88        2000
```

Prediction Confusion Matrix:

```
-----
              Predicted:
              positive  negative
Actual: positive      857      124
       negative      114      905
```

Newer Supervised Deep Learning Models

In [50]:

```
import gensim
import keras
from keras.models import Sequential
from keras.layers import Dropout, Activation, Dense
from sklearn.preprocessing import LabelEncoder

import spacy
import nltk
from nltk.tokenize.toktok import ToktokTokenizer

tokenizer = ToktokTokenizer()

nlp = spacy.load('en_core_web_sm')
```

Prediction class label encoding

In [23]:

```

le = LabelEncoder()
num_classes=2
# tokenize train reviews & encode train labels
tokenized_train = [tokenizer.tokenize(text)
                    for text in norm_train_reviews]
y_tr = le.fit_transform(train_sentiments)
y_train = keras.utils.np_utils.to_categorical(y_tr, num_classes)
# tokenize test reviews & encode test labels
tokenized_test = [tokenizer.tokenize(text)
                  for text in norm_test_reviews]
y_ts = le.fit_transform(test_sentiments)
y_test = keras.utils.np_utils.to_categorical(y_ts, num_classes)

```

In [24]:

```

# print class label encoding map and encoded labels
print('Sentiment class label map:', dict(zip(le.classes_, le.transform(le.clas
print('Sample test label transformation:\n'+ '-'*35,
      '\nActual Labels:', test_sentiments[:3], '\nEncoded Labels:', y_ts[:3],
      '\nOne hot encoded Labels:\n', y_test[:3])

```

Sentiment class label map: {'negative': 0, 'positive': 1}

Sample test label transformation:

Actual Labels: ['negative' 'negative' 'negative']

Encoded Labels: [0 0 0]

One hot encoded Labels:

[[1. 0.]

[1. 0.]

[1. 0.]]

Feature Engineering with word embeddings

In [26]:

```

# build word2vec model
w2v_num_features = 500
w2v_model = gensim.models.Word2Vec(tokenized_train, vector_size=w2v_num_featu
                                   min_count=10, sample=1e-3)

```

```
In [54]: def averaged_word2vec_vectorizer(corpus, model, num_features):
          vocabulary = set(model.wv.index_to_key)

          def average_word_vectors(words, model, vocabulary, num_features):
              feature_vector = np.zeros((num_features,), dtype="float64")
              nwords = 0.

              for word in words:
                  if word in vocabulary:
                      nwords = nwords + 1.
                      feature_vector = np.add(feature_vector, model.wv[word])
              if nwords:
                  feature_vector = np.divide(feature_vector, nwords)

              return feature_vector

          features = [average_word_vectors(tokenized_sentence, model, vocabulary, num_features)
                      for tokenized_sentence in corpus]
          return np.array(features)
```

```
In [56]: # generate averaged word vector features from word2vec model
          avg_wv_train_features = averaged_word2vec_vectorizer(corpus=tokenized_train,
                                                                num_features=500)
          avg_wv_test_features = averaged_word2vec_vectorizer(corpus=tokenized_test, model=model,
                                                             num_features=500)
```

```
In [57]: # feature engineering with GloVe model
          train_nlp = [nlp(item) for item in norm_train_reviews]
          train_glove_features = np.array([item.vector for item in train_nlp])

          test_nlp = [nlp(item) for item in norm_test_reviews]
          test_glove_features = np.array([item.vector for item in test_nlp])
```

```
In [58]: print('Word2Vec model:> Train features shape:', avg_wv_train_features.shape,
          print('GloVe model:> Train features shape:', train_glove_features.shape, 'Test features shape:', test_glove_features.shape)
```

```
Word2Vec model:> Train features shape: (5000, 500) Test features shape: (2000, 500)
GloVe model:> Train features shape: (5000, 96) Test features shape: (2000, 96)
```

Modeling with deep neural networks

Building Deep neural network architecture

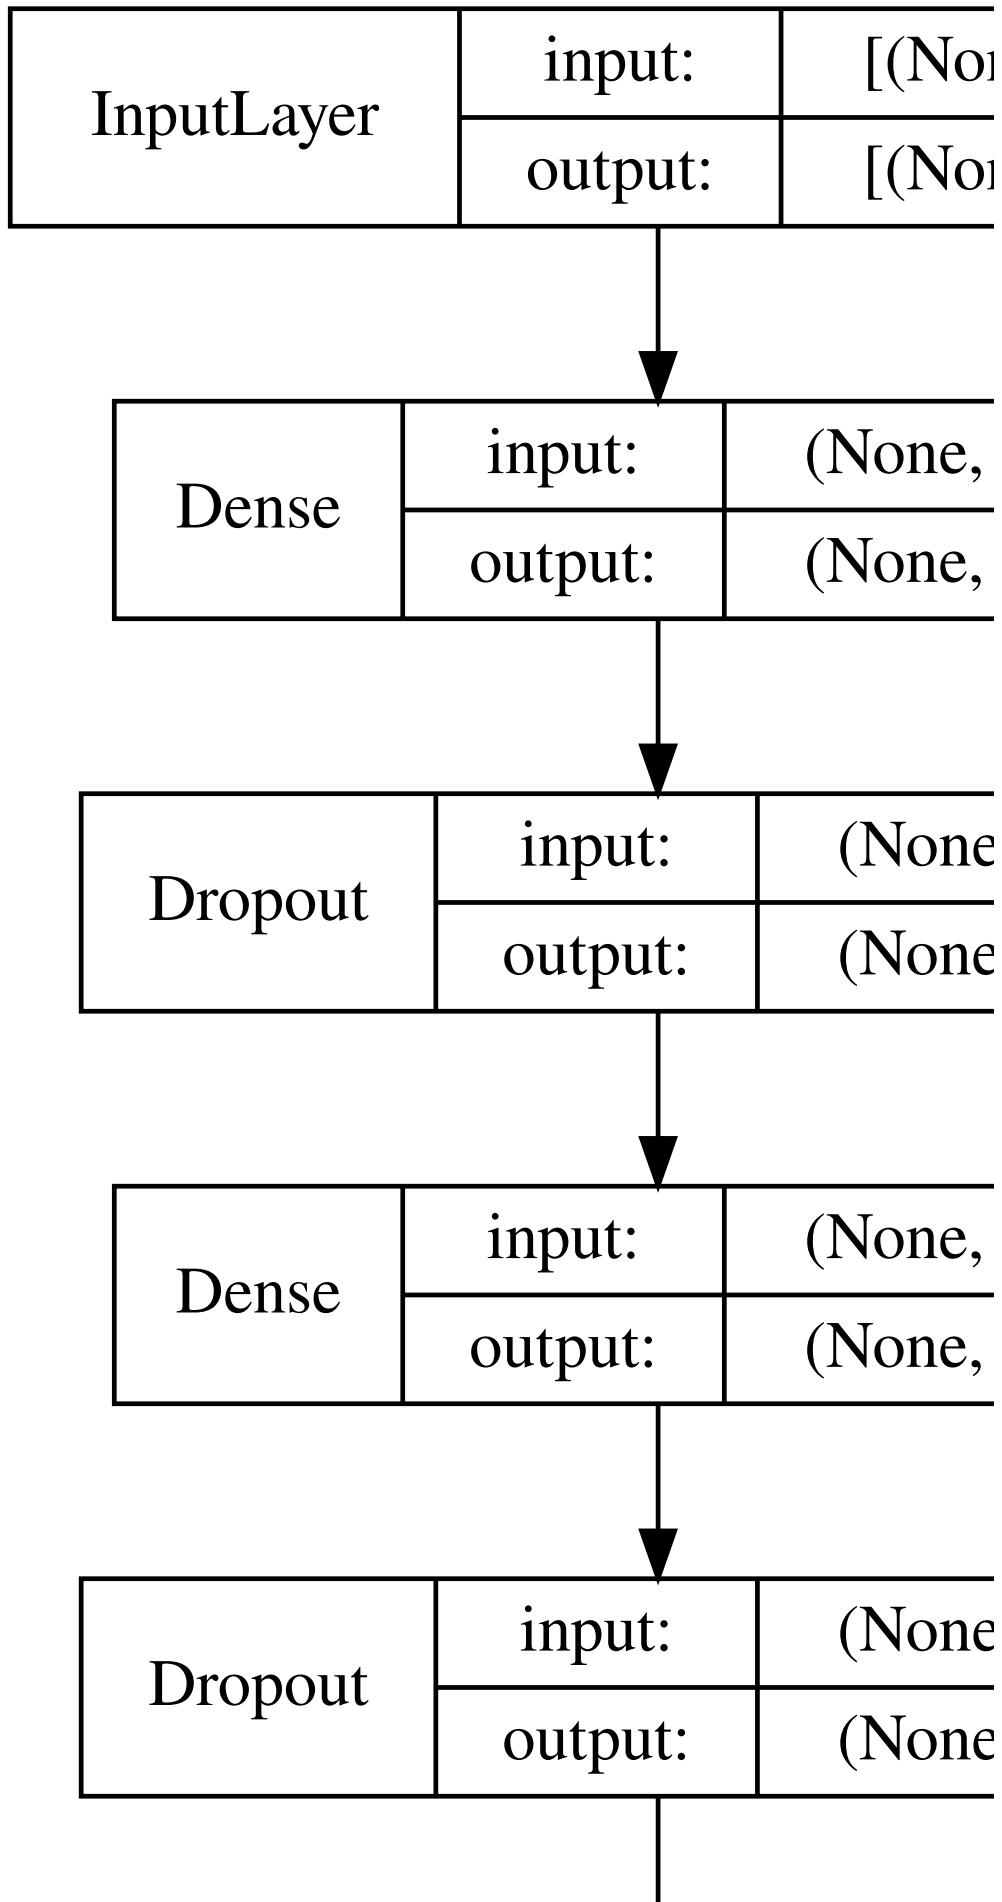
```
In [59]: def construct_deepnn_architecture(num_input_features):  
    dnn_model = Sequential()  
    dnn_model.add(Dense(512, activation='relu', input_shape=(num_input_features,)))  
    dnn_model.add(Dropout(0.2))  
    dnn_model.add(Dense(512, activation='relu'))  
    dnn_model.add(Dropout(0.2))  
    dnn_model.add(Dense(512, activation='relu'))  
    dnn_model.add(Dropout(0.2))  
    dnn_model.add(Dense(2))  
    dnn_model.add(Activation('softmax'))  
  
    dnn_model.compile(loss='categorical_crossentropy', optimizer='adam',  
                      metrics=['accuracy'])  
    return dnn_model
```

```
In [60]: w2v_dnn = construct_deepnn_architecture(num_input_features=500)
```

Visualize sample deep architecture

```
In [61]: from IPython.display import SVG  
    from keras.utils.vis_utils import model_to_dot  
  
    SVG(model_to_dot(w2v_dnn, show_shapes=True, show_layer_names=False,  
                     rankdir='TB').create(prog='dot', format='svg'))
```


Out[61]:





Model Training, Prediction and Performance Evaluation

```
In [62]: batch_size = 100
w2v_dnn.fit(avg_wv_train_features, y_train, epochs=5, batch_size=batch_size,
            shuffle=True, validation_split=0.1, verbose=1)
```

Epoch 1/5
45/45 [=====] - 1s 9ms/step - loss: 0.5115 - accuracy: 0.7516 - val_loss: 0.4990 - val_accuracy: 0.7640
Epoch 2/5
45/45 [=====] - 0s 5ms/step - loss: 0.4584 - accuracy: 0.7884 - val_loss: 0.4626 - val_accuracy: 0.7840
Epoch 3/5
45/45 [=====] - 0s 4ms/step - loss: 0.4469 - accuracy: 0.7913 - val_loss: 0.4565 - val_accuracy: 0.7840
Epoch 4/5
45/45 [=====] - 0s 4ms/step - loss: 0.4452 - accuracy: 0.7920 - val_loss: 0.4399 - val_accuracy: 0.7820
Epoch 5/5
45/45 [=====] - 0s 4ms/step - loss: 0.4356 - accuracy: 0.7938 - val_loss: 0.4483 - val_accuracy: 0.7820
Out[62]: <keras.callbacks.History at 0x7f529bcc7e10>

```
In [63]: #y_pred = w2v_dnn.predict_classes(avg_wv_test_features)
y_pred = w2v_dnn.predict(avg_wv_test_features)
y_classes = np.argmax(y_pred,axis=1)
predictions = le.inverse_transform(y_classes)
```

```
In [64]: import pkg_resources
pkg_resources.get_distribution('gensim').version
```

Out[64]: '4.1.2'

```
In [65]: meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_
                                             classes=['positive', 'negative'])
```

Model Performance metrics:

```
-----
Accuracy: 0.799
Precision: 0.8019
Recall: 0.799
F1 Score: 0.7988
```

Model Classification report:

```
-----
              precision    recall  f1-score   support

   positive       0.77       0.84       0.80        981
   negative       0.83       0.76       0.79       1019

   accuracy                   0.80        2000
  macro avg       0.80       0.80       0.80        2000
 weighted avg       0.80       0.80       0.80        2000
```

Prediction Confusion Matrix:

```
-----
              Predicted:
              positive  negative
Actual: positive      826      155
       negative      247      772
```

```
In [66]: glove_dnn = construct_deepnn_architecture(num_input_features=96)
```

```
In [67]: batch_size = 100
         glove_dnn.fit(train_glove_features, y_train, epochs=5, batch_size=batch_size,
                       shuffle=True, validation_split=0.1, verbose=1)
```

Epoch 1/5

45/45 [=====] - 1s 7ms/step - loss: 0.6740 - accuracy: 0.5836 - val_loss: 0.6558 - val_accuracy: 0.6140

Epoch 2/5

45/45 [=====] - 0s 5ms/step - loss: 0.6528 - accuracy: 0.6167 - val_loss: 0.6437 - val_accuracy: 0.6180

Epoch 3/5

45/45 [=====] - 0s 6ms/step - loss: 0.6410 - accuracy: 0.6293 - val_loss: 0.6375 - val_accuracy: 0.6460

Epoch 4/5

45/45 [=====] - 0s 6ms/step - loss: 0.6300 - accuracy: 0.6569 - val_loss: 0.6730 - val_accuracy: 0.5980

Epoch 5/5

45/45 [=====] - 0s 5ms/step - loss: 0.6393 - accuracy: 0.6356 - val_loss: 0.6387 - val_accuracy: 0.6300

```
Out[67]: <keras.callbacks.History at 0x7f52ab69e710>
```

```
In [68]: #y_pred = glove_dnn.predict_classes(test_glove_features)
y_pred = glove_dnn.predict(test_glove_features)
y_classes = np.argmax(y_pred,axis=1)
predictions = le.inverse_transform(y_classes)
```

```
In [69]: meu.display_model_performance_metrics(true_labels=test_sentiments, predicted_
                                             classes=['positive', 'negative'])
```

Model Performance metrics:

Accuracy: 0.6265

Precision: 0.6488

Recall: 0.6265

F1 Score: 0.6149

Model Classification report:

	precision	recall	f1-score	support
positive	0.59	0.81	0.68	981
negative	0.71	0.45	0.55	1019
accuracy			0.63	2000
macro avg	0.65	0.63	0.62	2000
weighted avg	0.65	0.63	0.61	2000

Prediction Confusion Matrix:

	Predicted:	
	positive	negative
Actual: positive	791	190
negative	557	462