

```
1 # ML1020 - Assignment 2 (DAG for wordcount)
2
3 """Example Airflow DAG that creates a Cloud Dataproc
4 cluster, runs the Hadoop
5 wordcount example, and deletes the cluster.
6
7 This DAG relies on three Airflow variables
8 https://airflow.apache.org/docs/apache-airflow/stable/
9 concepts/variables.html
10 * gcp_project - Google Cloud Project to use for the
11 Cloud Dataproc cluster.
12 * gce_zone - Google Compute Engine zone where Cloud
13 Dataproc cluster should be
14 created.
15 * gcs_bucket - Google Cloud Storage bucket to use for
16 result of Hadoop job.
17 See https://cloud.google.com/storage/docs/creating-
18 buckets for creating a
19 bucket.
20 """
21
22 import datetime
23 import os
24
25 from airflow import models
26 from airflow.providers.google.cloud.operators import
27 dataproc
28 from airflow.utils import trigger_rule
29
30 # Output file for Cloud Dataproc job.
31 # If you are running Airflow in more than one time zone
32 # see https://airflow.apache.org/docs/apache-airflow/
33 stable/timezone.html
34 # for best practices
35 output_file = os.path.join(
36     models.Variable.get('gcs_bucket'), 'wordcount',
37     datetime.datetime.now().strftime('%Y%m%d-%H%M%S'
38 )) + os.sep
```

```
30 # Path to Hadoop wordcount example available on every
    Dataproc cluster.
31 WORDCOUNT_JAR = (
32     'file:///usr/lib/hadoop-mapreduce/hadoop-mapreduce-
    examples.jar'
33 )
34 # Arguments to pass to Cloud Dataproc job.
35 #input_file = 'gs://pub/shakespeare/rose.txt'
36 input_file = 'gs://ml1020-bucket/asn1/*.txt'
37 wordcount_args = ['wordcount', input_file, output_file]
38
39 HADOOP_JOB = {
40     "reference": {"project_id": models.Variable.get('
    gcp_project')},
41     "placement": {"cluster_name": 'composer-hadoop-
    tutorial-cluster-{{ ds_nodash }}'},
42     "hadoop_job": {
43         "main_jar_file_uri": WORDCOUNT_JAR,
44         "args": wordcount_args,
45     },
46 }
47
48 CLUSTER_CONFIG = {
49     "master_config": {
50         "num_instances": 1,
51         "machine_type_uri": "n1-standard-2"
52     },
53     "worker_config": {
54         "num_instances": 2,
55         "machine_type_uri": "n1-standard-2"
56     },
57 }
58
59 yesterday = datetime.datetime.combine(
60     datetime.datetime.today() - datetime.timedelta(1),
61     datetime.datetime.min.time())
62
63 default_dag_args = {
```

```

64     # Setting start date as yesterday starts the DAG
        immediately when it is
65     # detected in the Cloud Storage bucket.
66     'start_date': yesterday,
67     # To email on failure or retry set 'email' arg to
        your email and enable
68     # emailing here.
69     'email_on_failure': False,
70     'email_on_retry': False,
71     # If a task fails, retry it once after waiting at
        least 5 minutes
72     'retries': 1,
73     'retry_delay': datetime.timedelta(minutes=5),
74     'project_id': models.Variable.get('gcp_project'),
75     'location': models.Variable.get('gce_region'),
76
77 }
78
79
80 with models.DAG(
81     'wordcount_assignment',
82     # Continue to run DAG once per day
83     schedule_interval=datetime.timedelta(days=1),
84     default_args=default_dag_args) as dag:
85
86     # Create a Cloud Dataproc cluster.
87     create_dataproc_cluster = dataproc.
        DataprocCreateClusterOperator(
88         task_id='create_dataproc_cluster',
89         # Give the cluster a unique name by appending
            the date scheduled.
90         # See https://airflow.apache.org/docs/apache-
            airflow/stable/macros-ref.html
91         cluster_name='composer-hadoop-tutorial-cluster
            -{{ ds_nodash }}',
92         cluster_config=CLUSTER_CONFIG,
93         region=models.Variable.get('gce_region'))
94

```

```
95     # Run the Hadoop wordcount example installed on  
    the Cloud Dataproc cluster  
96     # master node.  
97     run_wordcount = dataproc.DataprocSubmitJobOperator  
    (  
98         task_id='run_dataproc_hadoop',  
99         job=HADOOP_JOB)  
100  
101     # Delete Cloud Dataproc cluster.  
102     delete_dataproc_cluster = dataproc.  
    DataprocDeleteClusterOperator(  
103         task_id='delete_dataproc_cluster',  
104         cluster_name='composer-hadoop-tutorial-cluster  
-{{ ds_nodash }}',  
105         region=models.Variable.get('gce_region'),  
106         # Setting trigger_rule to ALL_DONE causes the  
    cluster to be deleted  
107         # even if the Dataproc job fails.  
108         trigger_rule=trigger_rule.TriggerRule.ALL_DONE  
    )  
109  
110     # Define DAG dependencies.  
111     create_dataproc_cluster >> run_wordcount >>  
    delete_dataproc_cluster
```